

---

# Explicating Obscure Text and Philosophizing via Unsupervised Text Style Transferring

---

Bin Li<sup>1</sup>

## Abstract

Writing is a serious endeavor. It takes years of training for one to become fully proficient. When humans employ writing as vehicle for ideas, the style of writing evolves accordingly, becoming more specialized. Certain styles of writing in some humanities fields have been criticized for practicing deliberate obscurantism to create a veneer of academic prestige. For this project, we create two datasets of philosophical writings in pedantic and explicative style respectively, and explore the possibility of using a neural text style transferring model to translate writings from one style to another. Particularly, we implement the Style Transfer via Paraphrasing (STRAP) method for this problem and also propose a novel modification to this method by integrating it into a multi-task learning framework. Our results show that the STRAP model that uses multi-task learning objective shows performance regression compared to the unmodified method. We utilize the unmodified STRAP model for our original datasets and qualitatively evaluate the performance of the model.

## 1. Introduction

It all started when Martha Nussbaum, a philosophy professor at the University of Chicago, wrote a lengthy article on The New Republic panning her fellow scholar Judith Butler for bad writing, whose works, Nussbaum argues, are nothing more than sophistry and should be precluded from genuine philosophical discussion (Nussbaum, 1999). "Academese," language that is filled with unnecessary jargons and turgid verses commonly associated with some fields of humanities, has drawn criticism aplenty. A common response to the criticism from within the academy is that the specialized

lingo is a necessity driven by the complexity and abstractness of subject matters. However, Pinker (2014) reasons that the roots of academese are in the mix of academics' goal to share their knowledge with the readers and their fear of "being convicted of philosophical naïveté about his own enterprise." (Thomas & Turner, 2017)

Is the convoluted mannerism of academic writing an inseparable blend of style and content or some fancy sprinkles on the cake of knowledge? In this project, we dive deeper into this problem with the help of the state-of-the-art self-supervised language modeling methods. The goal of this project is to perform a domain-specific expertise language translation/style transfer. Specifically, given some text written in a pedantic, obscurantist way, we want to output text that is similar in content but written in an explicative fashion. Here *pedantic* is defined to be a pompous writing style of deliberate mannerism that inhibits further inquiry, while *explicative* as a writing style that aims to clearly explain the logic of statement and avoid confusion.

Particularly, we experiment with one of the recent unsupervised neural text style transferring model called STRAP, short for Style Transfer via Paraphrasing (Krishna et al., 2020) and propose to incorporate the STRAP model in multi-task learning framework to further improve transfer accuracy. We show that our multi-task design negatively impact the model's performance on transfer quality, likely due to the unstable adversarial objectives during multi-task training.

Hence we implement the original STRAP on Pedantic/Explicative dataset, and qualitatively examine the style transfer performance. We found that the text generated by our pedantic/explicative STRAP model exhibit many linguistic characteristics that specific to their styles, such as style-specific terminology and syntax. Yet the model also discard the original information and add extra information not presented in the original semantics.

---

<sup>1</sup>Department of Computer Science, Columbia University, New York, NY. Correspondence to: Bin Li <bl2899@columbia.edu>.

## 2. Related Works

### 2.1. Academic Obscurantism

As prefaced, the argument about the mannerism of academic writing, especially in humanities disciplines, has been ongoing for the past decades. One particular malicious form is called academic obscurantism. Obscurantism is defined as “the practice of deliberately presenting information in an imprecise, abstruse manner designed to limit further inquiry and understanding.” The obscurantist writings are those that employ florid rhetorical devices, such as uncommon words, complex sentence structures, and exclusive literary references, to fool their readers into thinking they are more significant than they actually are. The obscurantist writing forces its reader to grant that, if one does not understand what is going on, there must be something important going on (Nussbaum, 1999).

In humanities, there are certain schools of thought commonly associated with academic obscurantism; a highly influential one is the school of postmodernism. Postmodernist philosophers and some literary theorists under their influence are often accused of committing deliberate obscurantism. In reaction to the growing interests in postmodernist teachings in the higher education institutions and concern for their apparent lack of substance, Alan Sokal, a physics professor at New York University, wrote a nonsensical paper about the “Hermeneutics of Quantum Gravity” and successfully published it on a postmodern cultural studies journal. This controversy shakes up the humanities circle and raises the question about the rigor of some humanities research.

Learning to understand these obscurantist writings in humanities fields poses a unique machine learning challenge. As those obscure philosophy essays and literary analyses are arguable the pinnacle of literary complexity in natural language, solving the challenge will be a major leap forward for natural language understanding.

### 2.2. Text Style Transfer

The goal of this project bears similarities with the problems of text simplification and text style transfer, which are two of the staple problems in the literature of natural language processing. Recent advancement in neural machine translation has offered some intriguing solutions on the problem of text simplification by recasting it as a problem of monolingual translation. Models like OpenNMT (Klein et al., 2017) offer a paradigm for supervised approach to the problem while unsupervised solution like UNTS (Surya et al., 2019) were also proposed. Those models, however, still fail to provide a meaningful solution to the problem, as they often perform shallow transformation on the lexical level. Their failure could be hinting at some deeper problems in the formula-

tion of the tasks of text simplification and text style transfer, particularly, the problem with our definitions of text style and content, and the question about their separability.

According to Cao et al. (2020), the approaches to text style transfer problem can be broken down into three categories. The first one is Disentanglement, which assumes that style and content can be separated into two distinct feature spaces and the model’s task is to disentangle the style vector from the content of data. ControlledGen (Hu et al., 2017), for example, employs variational autoencoder (VAE) to perform the distanglement for text style transfer. Another camp is translation, which does not assume the separability of style and content but instead learn style-specific transformation. Some examples of that include StyleTransformer (Dai et al., 2019), which use transformer models with cyclic architecture to learn style and content transformation, and Style Transfer via Paraphrasing, i.e., STRAP (Krishna et al., 2020), which creates style-specific pseudo-parallel text using large decoder language model. The last one is Manipulation, which leverages on text statistics to derive text-manipulating solution for style transfer. An example of that is DeleteAndRetrieve (Li et al., 2018), which utilizes statistical methods to identify words of certain style and replace them with those of another style.

## 3. Methods

In this project, we focus on using and modifying STRAP model for the task of text style transfer. We experiment with a novel modification to the model by incorporating STRAP model in multi-task training framework and compare their performance before deciding which model to use for our research problem.

### 3.1. STRAP Model


STRAP is a model for transferring styles between corpora of text without the support of explicit parallel example. It works by creating pseudo-parallel, style-neural dataset for each style using a style-neural paraphraser, which is a decoder language model that is able to generate diverse and style-neural paraphrase given an input sentence. This model is constructed through the following three steps.

Step one, we will prepare the model by finetuning a style-neural paraphraser  $P_N$  to generate more diverse outputs. A paraphraser  $P$  is a function such that when given an input sentence  $[x_1, \dots, x_n]$  in some corpus in style A, denoted  $\mathbf{X}^A$ ,

$$P([x_1, \dots, x_n]) = [z_1, \dots, z_m],$$

where sentences  $[x_1, \dots, x_n]$  and  $[z_1, \dots, z_m]$  are not identical yet semantically similar. A diverse paraphraser will output  $[z_1, \dots, z_m]$  that is maximally different  $[x_1, \dots, x_n]$  with in term of syntax and vocabulary.

## Neural machine translation using decoder LM



## Preparing a diverse paraphraser




Figure 1. Step one of STRAP

To encourage the diversity of  $P_N$ , STRAP finetunes  $P_N$  on a syntactically and lexically diverse back-translation dataset created by aggressively filtering indistinctive data from PARANMT-50M (Wieting & Gimpel, 2018). Paraphraser is trained using  $L_{LM}$  language modeling loss for next word prediction, i.e., perplexity, which is defined as

$$Perplexity(x_1, \dots, x_n) = \prod_{i=1}^n p_M(x_i | x_0, \dots, x_{i-1})^{-1/n},$$

where  $p_M(x_n | x_0, \dots, x_{n-1})$  is the probability of a language model  $M$  will output  $x_n$  given input  $x_0, \dots, x_{n-1}$ .

There are many candidate methods for the paraphraser, and STRAP chooses GPT2-large (Radford et al., 2019) because of its ability to generate fluent and complex sentences (Krishna et al., 2020). To tune GPT2, a decoder language model, for the back-translation language modeling objective, STRAP uses an encoder-free seq2seq strategy described by Wolf et al. (2019), which concatenate input and expected output sentences together to feed into the decoder model. We also summarize step one of STRAP in Figure 1.

Step two involves creating and training on pseudo parallel dataset. It starts with two corpora in two different styles, denoted as  $A, B$ . We then use the style-neural diverse paraphraser  $P_N$  to paraphrase all sentences in  $A$ . Formally,  $P(A) = N_A$ , where  $N_A$  is a dataset that maps every sentence in  $A$  to a paraphrastic sentence in neural style  $N$ , i.e.,  $N_A$  is  $A$ 's pseudo-parallel dataset.

Next, we train a new paraphraser  $P_A$  with sentences in  $N_A$  as input and those in  $A$  as output for back-translation language modeling object. The resulting  $P_A$  is a function that maps from  $N$  to  $A$ , i.e.,

$$P_A : N \rightarrow A$$

Do the same for corpus  $B$  and we obtain

$$P_B : N \rightarrow B$$

. Step three will piece all the component together at inference time by using style-neural paraphraser  $P_N$  and style-specific paraphraser  $P_A, P_B$  to map sentence to a certain style. Formally, this can be expressed as:

$$P_B(P(A)) : A \rightarrow B$$

$$P_A(P(B)) : B \rightarrow A$$

Figure 2 summarizes steps two and three.

## 3.2. Multi-Task Model

Recently there has been a growing interest in utilizing multi-task learning paradigm for training language model for various application. At the heart of the multi-task learning is a tug-of-war of multiple learning objectives, in which model is learning to make trade-off that will satisfy multiple learning requirements that potentially lead to better generalization to the model's applications (Sener & Koltun, 2018).

Here we propose a novel modification that will incorporate a multi-task learning objective with the STRAP model in hope of producing a better performing model.

We start by recognizing that one desirable objective for a style-specific paraphraser  $P_A$  is to know how to separate  $A$  from  $B$ , which will potentially lead to generating a more unambiguous and stylistically distinctive sentence that follows that style of  $A$ . We achieve this by introducing a sequence classification task along with the language modeling task to the training of style-specific paraphraser in STRAP. The classification task is constructed by sampling some sentences from style-distinct corpora  $A, B$  and asking the model to identify the style of these sentences.

More formally, we can describe this multi-task framework by specifying a task pool  $T$ , that

$$T = T^{LM} \cup T^{CLS} = \{t_1^{LM}, \dots, t_n^{LM}\} \cup \{t_1^{CLS}, \dots, t_m^{CLS}\},$$

where  $T^{LM}$  is all the language modeling training examples and  $t_i^{LM} = (x_i^{LM}, y_i^{LM})$ , i.e.,  $t_i^{LM}$  is the tuple that contains  $i$ -th training input  $x_i^{LM}$  and label  $y_i^{LM}$  for the language modeling task; vice-versa for sequence classification task  $CLS$ .


For a language model  $M_\theta$  parameterized by  $\theta$ , the optimization at some step  $k$  for our proposed multi-task framework can be expressed as:

$$t_k = (x_k, y_k) \sim T$$

$$\theta_{k+1} = \arg \min_{\theta} \begin{cases} L_{CLS}(f_{CLS}(M_{\theta_k}(x_k)), y_k), & t_k \in T^{CLS} \\ L_{LM}(f_{LM}(M_{\theta_k}(x_k)), y_k), & t_k \in T^{LM} \end{cases}$$

where  $t_k$  is sampled from  $T$  under some distribution,  $f_{CLS}, f_{LM}$  is the task-specific prediction heads for the language model, and  $L_{CLS}, L_{LM}$  are the task-specific losses,

Training time:



Inference time:



Figure 2. Steps two and three of STRAP.

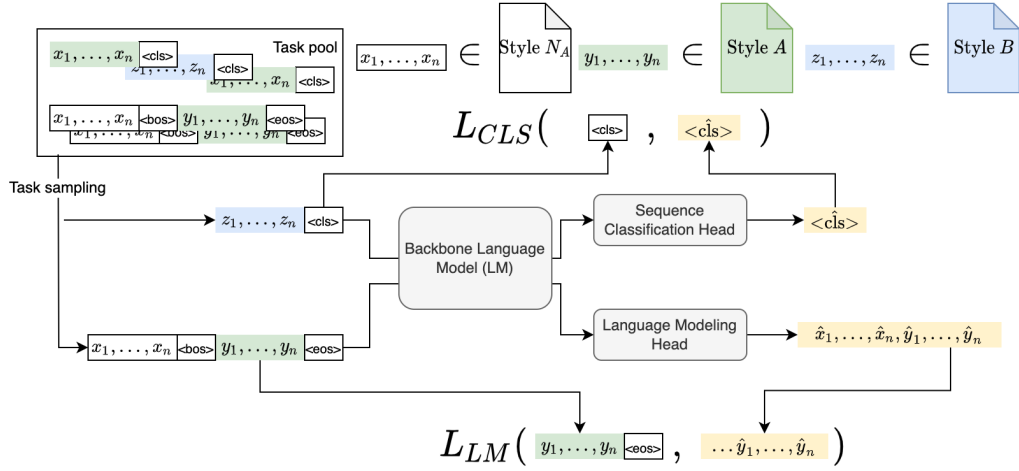


Figure 3. Multi-task paradigm for the STRAP model.

(cross-entropy for classification, perplexity for language modelling). Figure 3 provides a visual demonstration of this process.

### 3.3. Model Training

The models are trained on a single Nvidia V100 GPU using AdamW optimizer with the learning rate of  $2 \times 10^{-5}$ , weight decay rate of 0.001, and the batch size of 10. The training time ranges from 2 to 12 hours depending on the size of dataset. The models are evaluated using a single Nvidia 1080 Ti GPU and the evaluation time ranges from 1 to 3 hours.

## 4. Experiment

We divide our experiment into two stages. Firstly we test our proposed modified model against the original STRAP model along with other methods and choose the best one to perform the step two: style transferring between explicative and pedantic text.

### 4.1. Datasets

For the first stage of the experiment, we use the Shakespeare author imitation dataset (Xu et al., 2012), which contains 37k training sentences from two styles — William Shakespeare’s original plays, and their modernized versions. Shakespeare’s plays are written in Early Modern English, which has a significantly different lexical (e.g., thou instead of you) and syntactic distribution compared to modern En-

glish.

For the second stage, we generate two datasets: Explicative and Pedantic, by first identifying examples of pedantic and explicative writings in the fields of philosophy and literary theory. The set of pedantic text is composed of writings by philosopher and literary theorists who are commonly accused to be obscurantists. The list includes Jean Bau-drillard, Homi K. Bhabha, Judith Butler, Gilles Deleuze, Félix Guattari, Jacques Derrida, Michel Foucault, G. W. F. Hegel, Martin Heidegger, Julia Kristeva, Bruno Latour, and Emmanuel Levinas.

The explicative set is composed of relevant textbook materials and encyclopedia such as *Stanford Encyclopedia of Philosophy*, which is usually written in a clear, interpretative style in order to aid student’s understanding of the materials (see Appendix A for the full list of books used). After procuring the ebook files, we clean up the text to strip off any non-semantic information and extract sentences using regular expression. The size of the resultant set of pedantic writing is about 50k sentences while the explicative set has around 500k sentences, mainly because *Stanford Encyclopedia of Philosophy* is much longer than other documents. See Appendix B for sample writings from the dataset.

#### 4.2. Evaluation

We follow Krishna et al. (2020) and use three automatic evaluation methods for measuring text style transfer quality. These include: transfer accuracy (ACC), semantic similarity (SIM), and fluency (FL). ACC is measured by training a classifier to identify the style of the generated sentences. We use a fine-tuned RoBERTa-large (Liu et al., 2019) as the classifier for this task. For SIM, we use a subword embedding-based model proposed by Wieting et al. (2019), who argue that this method is a better alternative than the widely-used BLEU. For FL, we use a RoBERTa-large classifier trained on the CoLA corpus to identify the grammatical agreeability of sentences and use the identifying accuracy as FL score.

$J$  is an evaluation metric proposed by Krishna et al. (2020) to weight over three mentioned scoring methods on sentence-level to produce a single score that measures style transfer quality. It is defined as:

$$J(ACC, SIM, FL) = \sum_{x \in \mathbf{X}} \frac{ACC(x) \times SIM(x) \times FL(x)}{|\mathbf{X}|},$$

where  $x$  is a sentence in the test corpus  $\mathbf{X}$ .

Beside automatic evaluation metrics, we also make qualitative inspection of our results as human judgment is still irreplaceable in many task scenarios for natural language processing, and especially so for tasks like ours.

## 5. Results and Discussion

### 5.1. Shakespeare-Modern Style Transfer

The results of the automatic evaluation of style transfer quality are shown in Table 1. As the table shows, when putting STRAP in our proposed multi-task framework, we see an overall regression in performance. This likely indicates that the sequence classification task we add does not help the language modeling objective but rather impedes the optimization of the latter.

A quick inspection of the generated samples (Table 2) reveals many malign artifacts that the multi-task STRAP tends to produce; pathogenic patterns like starting a sentence with a special character (e.g. apostrophe, whitespace) or empty generation are common. While incompatibility could certainly be an important reason for the poor performance of the multi-task STRAP model, we should note that the size of the Shakespeare dataset is really modest, and it is hard to split up enough data for both language and classification objectives. Also notice that the one of the STRAP model’s generations for the Shakespeare style is exactly the same as the reference’s sentence for the modern style (“The other girl did not love me.”) This should raise some concerns about the possibility of test data leakage in STRAP’s methodology.

Nevertheless, STRAP still prevails over other models in our comparison, and our qualitative inspection also makes us convince its superiority compares to the multi-task variant. So we choose the original STRAP method for performance pedantic-explicative style transfer.

### 5.2. Pedantic-Explicative Style Transfer

We show a few samples of the generated outputs in Table 3. Judging qualitatively from the content of the style transferring, it seems that explicative style features shorter, clearer and more concise phrases. Yet the fidelity of the transferring of semantic meaning remains difficult to judge, primarily because the pedantic text in itself can be deliberately obscure.

We also test pedantic/explicative style transfer from more diverse style, using dataset like the collection of African American English tweets (Blodgett et al., 2016) or random tweets prepared by Krishna et al. (2020), and the results are shown in Table 4. Due to the lack of automatic evaluation, we need to rely again on qualitative human judgment. But a quick look at these samples unveils a few interesting insights about the pedantic/explicative paraphraser.

Firstly, the paraphraser can might discard or even distort the original meaning; as seen in the example of tweet “My little bro obviously wants me to stop supporting him financial”, where the explicative translation twists its meaning



Table 1. Comparison of automatic evaluation of text style transfer quality of the original STRAP and the multi-task STRAP along with other previously proposed methods. \* indicates results are obtained from Krishna et al. (2020).  $GM()$  is geometric mean.  $J()$  is an alternative metric discussed in Section 4.2. COPY is a lower-bound baseline, which passes input unmodified as output. REFERENCE is the gold-standard target. UNMT and DLSM are two state-of-the-art methods before STRAP.

Method	Shakespeare				
	Accuracy	Similarity	Fluency	$GM(ACC, SIM, FL)$	$J(ACC, SIM, FL)$
COPY*	9.6	67.1	79.1	37.1	7.2
REFERENCE*	90.4	100	79.1	89.4	70.5
UNMT*(Subramanian et al., 2018)	70.5	37.5	49.6	50.8	14.6
DLSM*(He et al., 2020)	71.1	43.5	49.4	53.5	16.3
STRAP ( $p = 0.0$ )*	71.7	<b>56.4</b>	<b>85.2</b>	<b>70.1</b>	<b>34.7</b>
STRAP ( $p = 0.9$ )*	<b>79.8</b>	47.6	71.7	64.8	27.5
Multi-Task STRAP (Ours)	55.2	42.1	64.1	53.0	15.5

to the opposite. This effect of “lost-in-translation” should be expected as the STRAP model need to paraphrase the text back and forth in order to map to the desired style.

Secondly, the paraphraser adds style-specific semantic information not presented in the original text. This is especially obvious when transferring from AAE or tweets to pedantic or explicative style. As AAE and tweets are extremely short, it is obvious that the lengthy pedantic/explicative translation is back-filled with extra information. Interestingly, the information being filled seems to exhibit style-specific characteristics. For example, the pedantic translation that says “essence of the being of others” might remind a knowledgeable reader of Heidegger’s writing (e.g. *Being and Time*), while explicative translation that says “which is the faculty by means of which I reason, that is to say, I think” seem to hearken back to the classical Descartes’s quote (“I think; therefore I am.”)

## 6. Conclusion and Future Work

In this project, we explore using unsupervised/self-supervised neural machine translation methods for transferring text style between pedantic and explicative text. Particularly, we experiment with the STRAP model and proposed a novel modification that reappropriate the STRAP model in multi-task learning framework, by concurrently training the STRAP model on language modeling and sequence classification task.

Our results shows that this multi-task design does not contribute positively to the model’s performance on transfer quality. So we choose to use the original STRAP as the style translator for our Pedantic/Explicative dataset. We qualitatively examine the style transfer performance on the Pedantic-Explicative and found that the GPT2-based STRAP model, while able to retain and instill many style-

specific characteristics in its generation, can discard the original and add unsolicited information to the translation.

Due to the limit of time and resource, we did not experiment with other style transferring models on the Pedantic-Explicative translation task. The future work can expand on the experiment by attempting other models such as cross-lingual language model (Conneau & Lample, 2019) and unsupervised neural machine translation (Subramanian et al., 2018), and recruit expert human subjects to provide evaluation.

Our project did not look too deep into the failure of the multi-task model. While it might be questionable how multi-task objectives can benefit unsupervised text style transferring, a closer look into the issues and a more detailed treatment as to why multi-task learning has led to performance regression in the STRAP model are still worth pursuing in future work.

## References

- Blodgett, S. L., Green, L., and O’Connor, B. Demographic dialectal variation in social media: A case study of African-American English. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 1119–1130, Austin, Texas, November 2016. Association for Computational Linguistics. doi: 10.18653/v1/D16-1120. URL <https://aclanthology.org/D16-1120>.
- Cao, Y., Shui, R., Pan, L., Kan, M.-Y., Liu, Z., and Chua, T.-S. Expertise style transfer: A new task towards better communication between experts and laymen. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 1061–1071, 2020.
- Conneau, A. and Lample, G. Cross-lingual language model

Table 2. Comparison of generated samples by the proposed multi-task method and the original STRAP model. Each cell shows generated sentence in the target style. REFERENCE is the gold standard translation. <s> denotes one whitespace character and <e> means empty generation. Bracketed, underlined text is comment and not part of the generation.

Shakespeare → Modern			Modern → Shakespeare		
REFERENCE	Multi-Task STRAP	STRAP	REFERENCE	Multi-Task STRAP	STRAP
And now you’ve changed ?	<s>change of mind?	And you’re go- ing to change your mind ?	And art thou changed ?	’re changing your mind?	And now change your mind ?
The other girl did not love me .	<s>other one didn’t really.	The other one didn’t .	The other did not so .	<s>was the other girl’s lover.	The other girl did not like me .
But come on , inconsistent young man , come with me .	<e>	Come , young man , come , come , come , <u>[repeating]</u>	But come , young waverer , come , go with me , In one respect I’ll thy assistant be , For this alliance may so happy prove To turn your households , rancor to pure love .	<s>is young, come with me.	But young man , come hither .

pretraining. *Advances in neural information processing systems*, 32, 2019.

Dai, N., Liang, J., Qiu, X., and Huang, X.-J. Style transformer: Unpaired text style transfer without disentangled latent representation. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 5997–6007, 2019.

He, J., Wang, X., Neubig, G., and Berg-Kirkpatrick, T. A probabilistic formulation of unsupervised text style transfer. *arXiv preprint arXiv:2002.03912*, 2020.

Hu, Z., Yang, Z., Liang, X., Salakhutdinov, R., and Xing, E. P. Toward controlled generation of text. In *International conference on machine learning*, pp. 1587–1596. PMLR, 2017.

Klein, G., Kim, Y., Deng, Y., Senellart, J., and Rush, A. M. Opennmt: Open-source toolkit for neural machine translation. In *Proceedings of ACL 2017, System Demonstrations*, pp. 67–72, 2017.

Krishna, K., Wieting, J., and Iyyer, M. Reformulating unsupervised style transfer as paraphrase generation. *arXiv preprint arXiv:2010.05700*, 2020.

Li, J., Jia, R., He, H., and Liang, P. Delete, retrieve, generate: a simple approach to sentiment and style transfer. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 1865–1874, 2018.

Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer, L., and Stoyanov, V. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.

Nussbaum, M. C. The professor of parody. *New Republic*, 220(8):37–45, 1999.

Pinker, S. Why academics stink at writing. *The chronicle of higher education*, 61(5), 2014.

Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I., et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.

Sener, O. and Koltun, V. Multi-task learning as multi-objective optimization. In Bengio, S., Wallach, H., Larochelle, H., Grauman, K., Cesa-Bianchi, N., and Garnett, R. (eds.), *Advances in Neural Information Processing Systems*, volume 31. Curran Associates,

Table 3. Samples of generated sentences from pedantic to explicative style.

Pedantic (original)	Explicative (transferred)
At this point, the value of representation in the common sense divides into two irreducible values in the para-sense: a distinctness which can only be obscure, the more obscure the more it is distinct; and a confusion-clarity which can only be confused.	The two irreducibly different values are, in fact, only masked by a more profound obscurity: the first value is the value of separation, while the second is that of confusion, that is, the negation of unity.
Since analytical cognition, the transformation of the given material into logical determinations, has shown itself to be a positing that immediately determines itself to be equally a presupposing, to be both in one, the logical element can appear on account of this presupposing to be in the subject matter as something already completed, just as because of the positing it can appear as the product of a merely subjective activity.	Analytical cognition shows us that material objects are produced by logical cognition, that is, they immediately become the objects of the two presuppositions, the subject and the predicate, thus appearing as the same object in both cases.

Inc., 2018. URL <https://proceedings.neurips.cc/paper/2018/file/432aca3a1e345e339f35a30c8f65edce-Paper.pdf>.

Subramanian, S., Lample, G., Smith, E. M., Denoyer, L., Ranzato, M., and Boureau, Y.-L. Multiple-attribute text style transfer. *arXiv preprint arXiv:1811.00552*, 2018.

Surya, S., Mishra, A., Laha, A., Jain, P., and Sankaranarayanan, K. Unsupervised neural text simplification. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 2058–2068, 2019.

Thomas, F.-N. and Turner, M. *Clear and Simple as the Truth*. Princeton University Press, 2017.

Wieting, J. and Gimpel, K. ParaNMT-50M: Pushing the limits of paraphrastic sentence embeddings with millions of machine translations. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 451–462, Melbourne, Australia, July 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1042. URL <https://aclanthology.org/P18-1042>.

Wieting, J., Berg-Kirkpatrick, T., Gimpel, K., and Neubig, G. Beyond BLEU: training neural machine translation with semantic similarity. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 4344–4355, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1427. URL <https://aclanthology.org/P19-1427>.

Wolf, T., Sanh, V., Chaumond, J., and Delangue, C. Transfertransfo: A transfer learning approach for neu-

ral network based conversational agents. *arXiv preprint arXiv:1901.08149*, 2019.

Xu, W., Ritter, A., Dolan, B., Grishman, R., and Cherry, C. Paraphrasing for style. In *Proceedings of COLING 2012*, pp. 2899–2914, Mumbai, India, December 2012. The COLING 2012 Organizing Committee. URL <https://aclanthology.org/C12-1177>.



Table 4. Samples of generated sentences from African American English (AAE) to pedantic and explicative style.

Original Style	Original Text	Pedantic (transferred)	Explicative (transferred)
AAE	My outfit tomorrow ??? ^_^	The next morning, I'm going to be wearing this outfit, and it will be the last thing I ever wear, tomorrow, the day that I die? Tomorrow, when I am no longer a human being, but a machine, a cog in the great machine that is the machine of death? To what end, to the end of my existence, will I wear this costume, this mask, in order to become, as it were, 'the other of the other'? Will it not be a masquerade, an act of masochism that will, at the same time, make me a part of this machine and, through it, transform me into a different person,	What tomorrow will look like for me? Will I be wearing the same outfit tomorrow as I was yesterday, or will I have changed into something else entirely, such as a pair of jeans or a t-shirt? What is the difference between my current outfit and tomorrow's outfit, and how can I tell which is which? These questions are
AAE	I Use To Think They Were Ugly .	I have thought them to be ugly, but I have also thought that they are beautiful, and that the beauty of the beautiful is that it is not the same as the ugly; that there is no beauty without ugliness, that beauty is always the ugliest of all.	I have thought them to be ugliness, and I have said that they are ugly, but I do not believe that there is any such thing as an ugly thing; I believe only in the existence of things that are not beautiful.
AAE	I Don't Know What I Wanna Do With My Life Smh	I don't know what life would be like for me, but I do know that I would not want to live it in a way that would deprive myself of what is most important to me: my relationship with God.	I don't know what kind of life I want to lead, but I do know that it is not the life of a philosopher, or of someone who has devoted his life to the study of philosophy, and that I have no desire to become a professional philosopher.
Tweets	This little human of mine is so pure & perfect.	I am such an idealized human being that I have attained such purity and perfection in my being, and in this purity, I am the absolute essence of the being of others.	I am such that I have a perfect and pure human nature, and I am endowed with a rational faculty, which is the faculty by means of which I reason, that is to say, I think.
Tweets	My little bro obviously wants me to stop supporting him financially	I will not be supporting my brother in any way, shape, or form, but I will be doing everything in my power to make sure that he does not become a victim of his own success.	I know that I am not supposed to be supporting my boyfriend financially, but I feel obligated to do so because he is my best friend and I love him, even if I don't know how to make ends meet. Am I morally required to pay him back? How can I be sure

**A. List of books in the Explicative and Pedantic datasets**

Author/s	Title	Style
Bertrand Russell	History of Western Philosophy	Explicative
Various Authors	The Philosophy Book	Explicative
Various Authors	Stanford Encyclopedia of Philosophy	Explicative
Jean Baudrillard	Simulacra and Simulation	Pedantic
Jean Baudrillard	Symbolic Exchange and Death	Pedantic
Homi K. Bhabha	The Location of Culture	Pedantic
Judith Butler	Gender Trouble	Pedantic
Judith Butler	Undoing Gender	Pedantic
Gilles Deleuze and Félix Guattari	Anti-Oedipus	Pedantic
Gilles Deleuze and Félix Guattari	A Thousand Plateaus	Pedantic
Michel Foucault	Discipline and Punish	Pedantic
Michel Foucault	Madness and Civilization	Pedantic
G. W. F. Hegel	The Phenomenology of Spirit	Pedantic
G. W. F. Hegel	The Science of Logic	Pedantic
Martin Heidegger	Being and Time	Pedantic
Martin Heidegger	Being and Truth	Pedantic
Martin Heidegger	Introduction to Metaphysics	Pedantic
Julia Kristeva	The Portable Kristeva	Pedantic
Bruno Latour	Laboratory Life	Pedantic
Bruno Latour	Reassembling The Social	Pedantic
Bruno Latour	We Have Never Been Modern	Pedantic
Emmanuel Levinas	Levinas Unhinged	Pedantic

## B. Samples of writing

Source	Text	Style
Stanford Encyclopedia of Philosophy	First, whether an individual has a complaint under Nozick's Lockean proviso will not be a matter of whether too little natural material has been left in common by others or too little material has been left for that individual to initially acquire.	Explicative
Stanford Encyclopedia of Philosophy	Part of Kant's aim in the "Critique of Teleological Judgment" is to clarify the relation of natural teleology to religion, and to argue in particular against "physicoteleology," that is, the attempt to use natural teleology to prove the existence of God.	Explicative
Russell - History of Western Philosophy	The only thought which philosophy brings with it to the contemplation of history is the simple conception of Reason; that Reason is the sovereign of the world; that the history of the world, therefore, presents us with a rational process.	Explicative
Kristeva - The Portable Kristeva	But at the same time and as a result, textual experience reaches the very foundation of the social—that which is exploited by sociality but which elaborates and can go beyond it, either destroying or transforming it.	Pedantic
Butler - Gender Trouble	If this analysis is right, then the juridical formation of language and politics that represents women as "the subject" of feminism is itself a discursive formation and effect of a given version of representational politics.	Pedantic
Deleuze - Difference and Repetition	Curie commented that it was useful but misleading to speak of dissymmetry in negative terms, as though it were the absence of symmetry, without inventing positive terms capable of designating the infinite number of operations with unmatched outcomes.	Pedantic
Heidegger - Being and Time	Confronted by a phenomenon such as conscience, what is ontologically and anthropologically inadequate about a free-floating framework of classified faculties of the soul or personal acts becomes painfully obvious.	Pedantic
Foucault - Discipline and Punish	And whatever theatrical elements it still retained were now downgraded, as if the functions of the penal ceremony were gradually ceasing to be understood, as if this rite that 'concluded the crime' was suspected of being in some undesirable way linked with it.	Pedantic
Hegel - The Phenomenology of Spirit	This renunciation, which the essence has already in itself completed, the acting self presents in existence and for its consciousness, and it replaces that former immediate actuality with the higher actuality, namely, that of himself.	Pedantic
Bhabha - The Location of Culture	Culture becomes as much an uncomfortable, disturbing practice of survival and supplementarity – between art and politics, past and present, the public and the private – as its resplendent being is a moment of pleasure, enlightenment or liberation.	Pedantic